

COT 6405
ANALYSIS OF ALGORITHMS

B-Trees

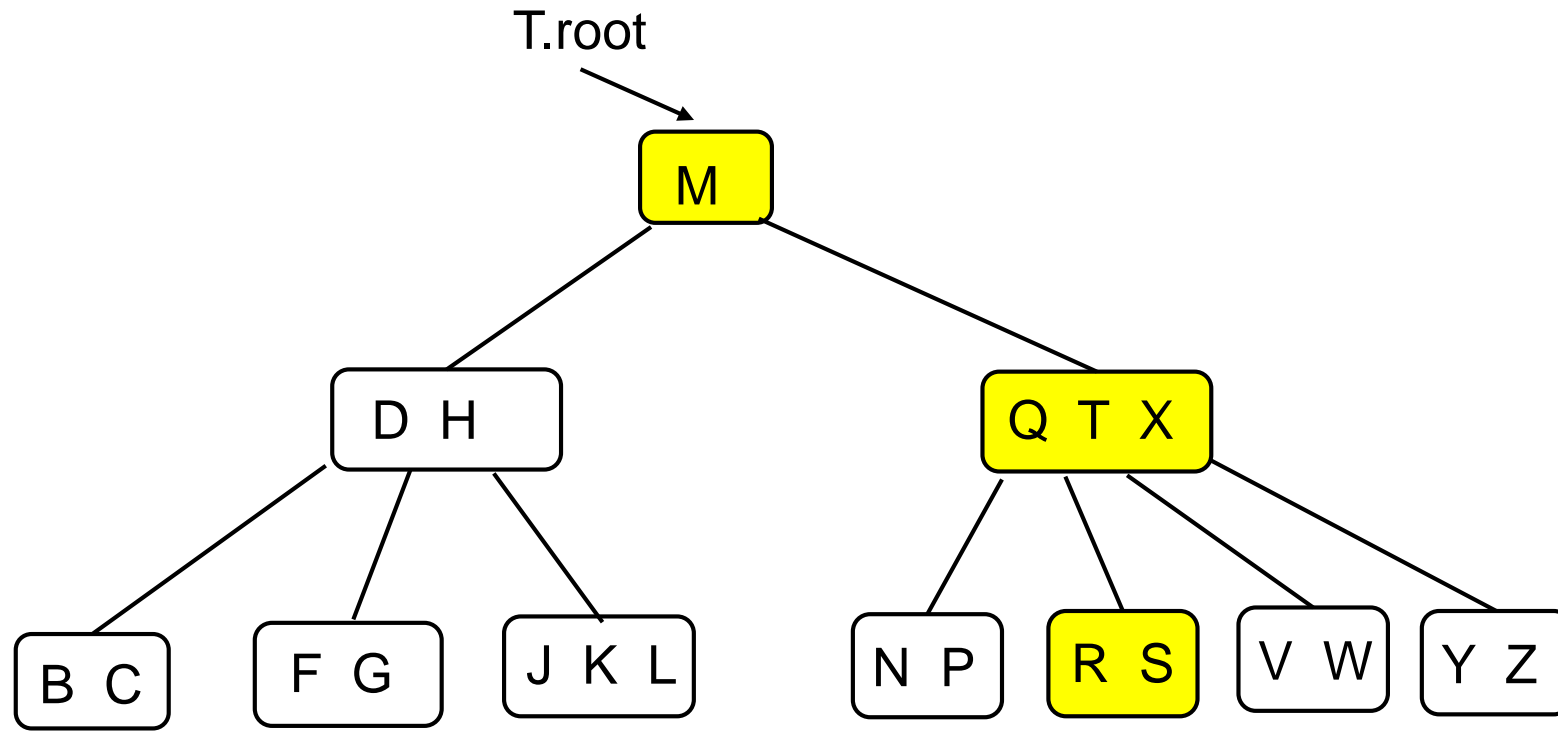
Computer & Electrical Engineering and Computer Science Department
Florida Atlantic University

B-trees

- Balanced search trees
- Work well on disks and other direct-access secondary storage devices
- Many database systems use B-trees or variants of B-trees to store information
- Efficient in minimizing disk I/O operations
- B-tree nodes may have from a few to thousands children
- B-trees have height $O(\lg n)$

Reference: *Introduction to Algorithms*, 3rd edition, by T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein, The MIT Press, 2009, chapter 18.

B-tree example



- keys are the consonants of English alphabet
- an internal node with $x.n$ keys has $x.n + 1$ children
- all leaves have the same depth
- the shaded nodes are examined in search for the letter R

Primary/Secondary storage

- The **primary memory** (main memory) consists of silicon memory chips
- **Secondary storage:** magnetic storage technology such as tapes or disks
- Disks are *cheaper* and have *higher capacity* than the main memory
- Disks are much *slower* than the main memory because they have moving mechanical parts

Primary/Secondary storage

- average access time for commodity disks is $\sim 8 - 11$ ms
 - access time for silicon memory is ~ 50 ns
- \Rightarrow access time for disks is over 5 order of magnitude slower!
- information divided in **pages** ($2^{11} - 2^{14}$ bytes)
 - each disk reads/writes one or more pages

B-trees

- in a typical B-tree application, the whole B-tree does not fit in the main memory
- copy pages from disk to the main memory (MM), then write back onto the disk the pages that have changed
- usually, B-tree algorithms keep only a constant number of nodes in the main memory (MM)

x = a pointer to some object

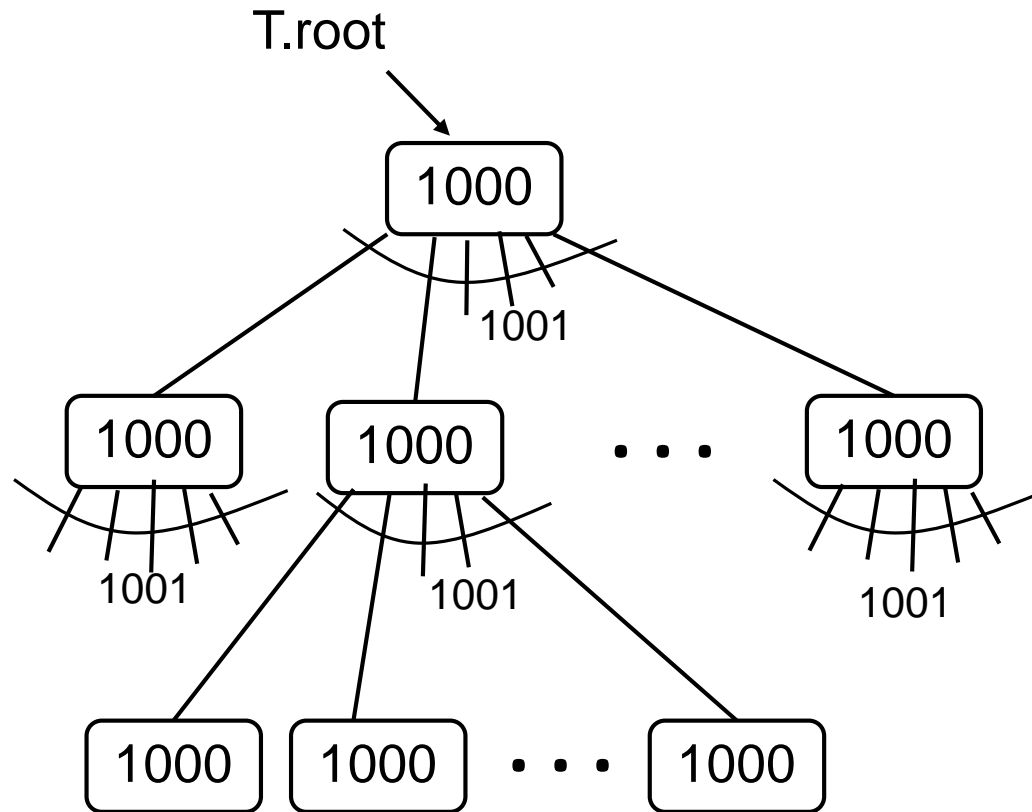
DISK-READ(x) // read object x in MM; no-op if
// x already in the MM

operations that access/modify attributes of x

DISK-WRITE(x) // omitted if no attributes of x changed

other operations that access but do not modify x

B-tree example: branching factor = 1001, height=2



1 node,
1000 keys

1001 nodes,
1,001,000 keys

1,002,001 nodes,
1,002,001,000 keys

- B-trees stored on disks, often have branching factors 50 ... 2000
- keep the root node permanently in the MM \Rightarrow find any key with at most two disk accesses

B-tree definition

A B-tree T is a rooted tree (where $T.root$ is the root) with the following properties:

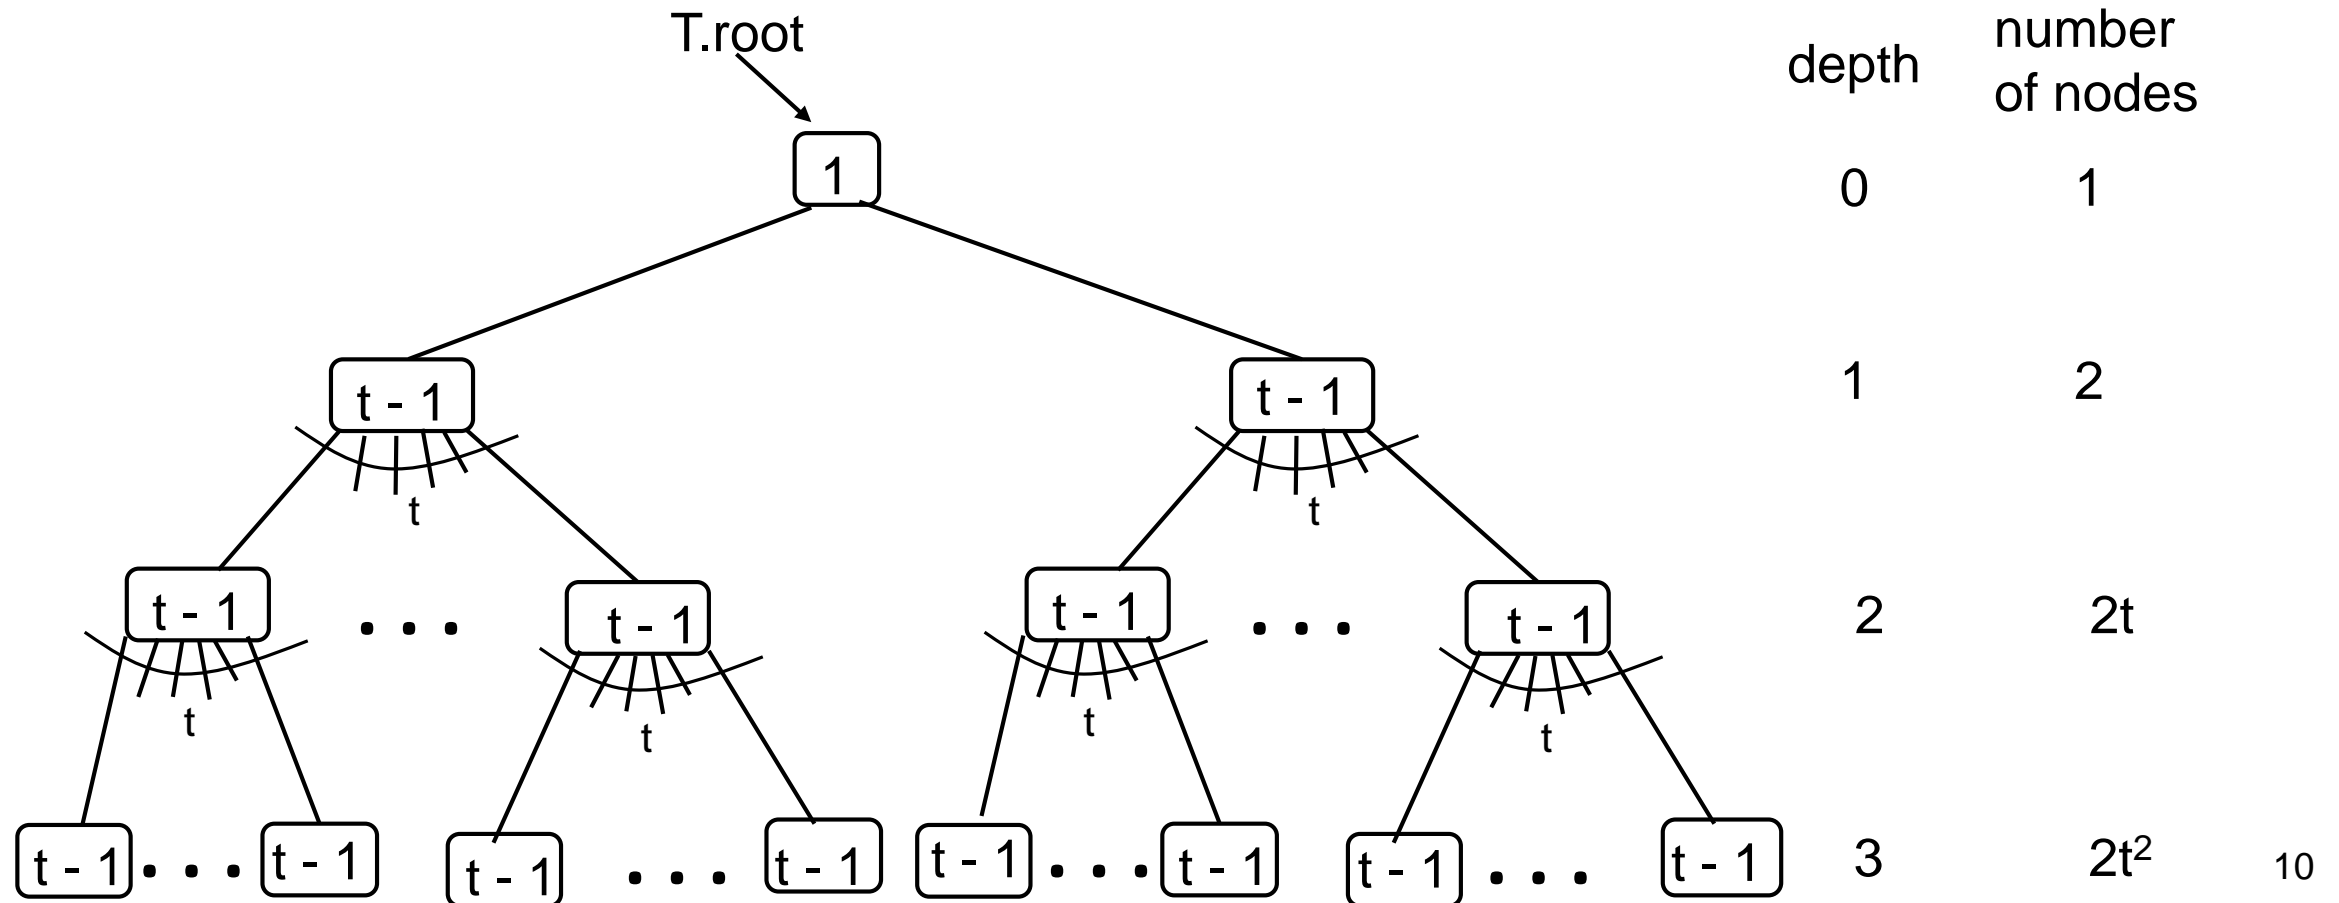
1. every node x has the following attributes
 - a. $x.n$ – the number of keys currently stored in x
 - b. the keys $x.key_1, x.key_2, \dots, x.key_{x.n}$ so that
$$x.key_1 \leq x.key_2 \leq \dots \leq x.key_{x.n}$$
 - c. $x.leaf$ – a boolean value which is TRUE if x is a leaf and FALSE if x is an internal node
2. each internal node x has $x.n+1$ pointers $x.c_1, x.c_2, \dots, x.c_{x.n+1}$ to its children; if x is a leaf then its c_i attributes are undefined
3. if k_i is any key stored in the subtree with root $x.c_i$ then:
$$k_1 \leq x.key_1 \leq k_2 \leq x.key_2 \leq \dots \leq x.key_{x.n} \leq k_{x.n+1}$$

B-tree definition, cont.

4. all the leaves have the same depth, which is the tree high h .
5. the B-tree has a **minimum degree** t (t is an integer $t \geq 2$):
 - every node other than the root must have $\geq t - 1$ keys and $\geq t$ children; if B-tree is nonempty, then the root has at least one key
 - every node has $\leq 2t - 1$ keys and $\leq 2t$ childrenA node is full is it has $2t - 1$ keys.

The height of a B-tree

Theorem: if $n \geq 1$, then for any n -key B-tree T of height h and minimum degree t ,

$$h \leq \log_t \frac{n+1}{2}$$


The height of a B-tree, cont.

$$n \geq 1 + (t-1) \sum_{i=1}^h 2t^{i-1} = 1 + 2(t-1) \sum_{i=1}^h t^{i-1}$$

$$= 1 + 2(t-1) \frac{t^h - 1}{t - 1} = 2t^h - 1$$

$$t^h \leq \frac{n+1}{2}$$

$$h \leq \log_t \frac{n+1}{2}$$

$$h = O(\log_t n)$$

$$h = O(\lg n)$$

Basic operations of B-trees

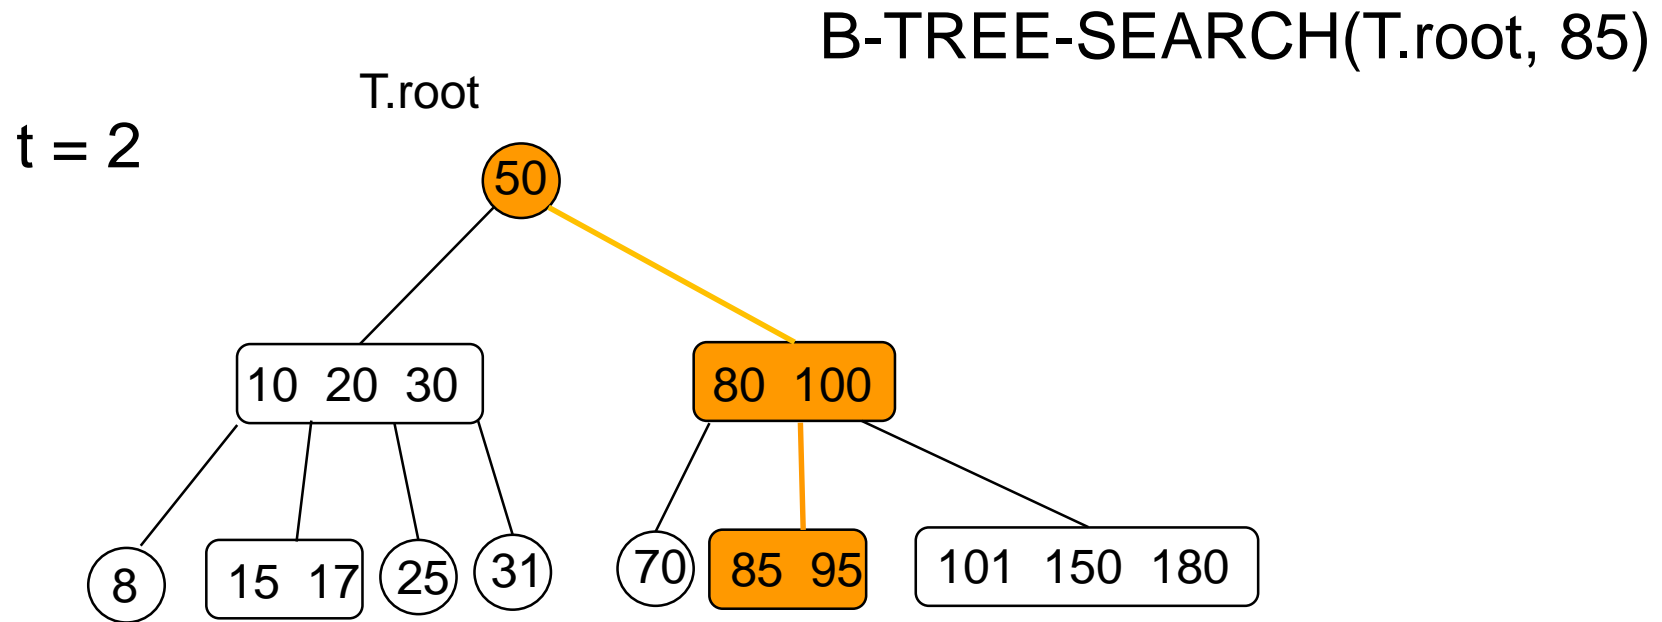
- B-TREE-SEARCH
- B-TREE-CREATE
- B-TREE-INSERT
- B-TREE-DELETE

Conventions

- The root of the B-tree is always in the main memory
 - no need to call DISK-READ for the root
 - we do need to call DISK-WRITE when the root is changed
- All nodes passed as parameters must have already had a DISK-READ operation performed on them
- The procedures are “**one-pass**” **algorithms** that proceed downward from the root, without having to back up

Searching a B-tree

- At each node make a $(x.n + 1)$ – way branching decision



Search operation

B-TREE-SEARCH(x, k)

$i = 1$

while $i \leq x.n$ and $k > x.key_i$

$i = i + 1$

if $i \leq x.n$ and $k == x.key_i$


return (x,i)

elseif $x.leaf == \text{TRUE}$

return NIL

else DISK-READ($x.c_i$)

return B-TREE-SEARCH($x.c_i$, k)



Initial call: B-TREE-SEARCH(T.root, k)

RT = $O(t \cdot \log_t n)$

- while loop takes $O(t)$
- number of recursive calls is $O(h) = O(\log_t n)$

Creating an empty B-tree

- Creates an empty root node

B-TREE-CREATE(T)

$x = \text{ALLOCATE-NODE}()$

$x.\text{leaf} = \text{TRUE}$

$x.n = 0$

$\text{DISK-WRITE}(x)$

$T.\text{root} = x$



$RT = O(1)$